Cache Bandwidth Allocation for P2P File-Sharing Systems to Minimize Inter-ISP Traffic

Valentino Pacifici, Frank Lehrieder, and György Dán

Abstract—**Many Internet service providers (ISPs) have deployed peer-to-peer (P2P) caches in their networks in order to decrease costly inter-ISP traffic. A P2P cache stores parts of the most popular contents locally, and if possible serves the requests of local peers to decrease the inter-ISP traffic. Traditionally, P2P cache resource management focuses on managing the storage resource of the cache so as to maximize the inter-ISP traffic savings. In this paper, we show that when there are many overlays competing for the upload bandwidth of a P2P cache, then in order to maximize the inter-ISP traffic savings, the cache's upload bandwidth should be actively allocated among the overlays. We formulate the problem of P2P cache bandwidth allocation as a Markov decision process and propose three approximations to the optimal cache bandwidth allocation policy. We use extensive simulations and experiments to evaluate the performance of the proposed policies, and show that the bandwidth allocation policy that prioritizes swarms with a small ratio of local peers to all peers in the swarm can improve the inter-ISP traffic savings in BitTorrent-like P2P systems by up to 30%–60%.**

Index Terms—**Caching, inter-ISP traffic management, Markov decision process, peer-to-peer content distribution.**

I. INTRODUCTION

THE NUMBER of peer-to-peer (P2P) applications has increased significantly in recent years, and so has the amount of Internet traffic generated by P2P applications. P2P traffic accounts for up to 70% of the total network traffic, depending on geographical location [1], and is a significant source of inter-Internet service provider (ISP) traffic. Inter-ISP traffic can be a source of revenue for tier-1 ISPs, but it is a source of transit traffic costs for ISPs at the lower levels of the ISP hierarchy, e.g., for tier-2 and tier-3 ISPs. Some ISPs have attempted to limit their costs due to P2P applications by throttling P2P traffic [2]. Nevertheless, the users of P2P applications constitute a significant share of the ISPs' customer base, and hence a solution that negatively affects the performance of P2P applications can result in a decrease of an ISP's revenues on the long term.

Recent research efforts have tried to decrease the amount of inter-ISP P2P traffic by introducing locality awareness in the neighbor-selection policies of popular P2P applications, like BitTorrent [3]–[6]. Locality information can be provided by the ISPs [3]–[5] or can be obtained via measurements [6] and is used to prioritize nearby peers to distant ones when exchanging data. Through exchanging data primarily with nearby peers, a P2P application can improve the locality of its traffic and, hence, can decrease inter-ISP traffic. Nevertheless, locality-aware neighbor selection can deteriorate the performance and the robustness of a P2P application [7].

To address the problem of increased inter-ISP traffic, many ISPs have deployed P2P caches [8], [9]. P2P caches, similar to Web proxy caches, decrease the amount of inter-ISP traffic by storing the most popular contents in the ISP's own network, so that they do not have to be downloaded from peers in other ISPs' networks. According to measurement studies, 30%–80% of P2P traffic is cacheable [10], [11]. Nevertheless, the actual efficiency of a cache depends on two main factors: first, the amount of storage, which determines the share of the contents that can be kept in cache; second, the available bandwidth of the cache, which determines the rate at which data can be served by the cache, if the data are in storage.

The goal of cache storage management is to maximize the probability that data are found in the cache when requested. The algorithms for cache storage management, called cache eviction policies, in the case of P2P caches differ significantly from those in the case of Web proxy caching. Web objects are typically small, and consequently eviction policies can replace entire contents at once [12]. Objects in P2P systems are nevertheless typically too big to be replaced at once, so that eviction policies for P2P caches have to allow partial caching of contents [10], [11]. By allowing partial caching, P2P eviction policies can achieve within 10%–20% of the optimal offline eviction policy [10], [11].

The impact of the cache bandwidth and its management has received little attention, even though cache bandwidth can be costly, as caches are often priced based on their bandwidth [8], [9]. In the case of Web proxy caching, bandwidth management is not necessary because the incoming inter-ISP traffic saving equals the amount of data served from the cache. In the case of a P2P cache, the inter-ISP traffic saving is, however, not only determined by how much data the cache serves but also by the swarm to which the data are served [13], [14]. As an example, consider two swarms, A and B, spread over two ISPs, 1 and 2. Swarm A has 1 peer in ISP 1 and 9 peers in ISP 2, while swarm B has 10 peers in ISP 1 and none in ISP 2. The 10 peers of swarm B would arguably use most

of the bandwidth of ISP 1's P2P cache, even though they do not generate any inter-ISP traffic. This example suggests that bandwidth allocation may be beneficial, but it is unclear when this would be the case and how bandwidth should be allocated.

The fundamental question we address in this paper is whether given a limited amount of P2P cache bandwidth, the bandwidth can be actively managed such as to minimize the amount of inter-ISP traffic. We make three important contributions to answer this question. First, we provide a mathematical formulation of the cache bandwidth allocation problem and show the existence of a stationary optimal policy. Second, we use the proposed mathematical model and insights from [13] and [14] to derive three allocation policies to approximate the optimal policy. Third, through simulations and through experiments on Planet-lab, we show that by actively allocating the upload bandwidth between different overlays the inter-ISP traffic, savings due to P2P caches can be improved significantly. We identify the heterogeneity of the ratio of local peers in the swarms as the key factor that determines the potential traffic savings.

The rest of the paper is organized as follows. In Section II, we review the related work. In Section III, we model the system and its evolution using a Markov jump process. In Section IV, we formulate the problem of cache bandwidth allocation and show the existence of an optimal cache bandwidth allocation policy. Section V describes three policies to approximate the optimal policy. In Section VI, we use simulation and experiment results to quantify the potential of the proposed bandwidth allocation policies and to provide insight into the characteristics of an optimal policy. Section VII concludes the paper.

II. RELATED WORK

The solutions for ISP-friendly P2P application design proposed in the literature fall into three main categories: peer-driven, ISP-driven, and caching [15]. Peer-driven solutions adapt the neighbor selection strategy of the peers by relying on measurements of latency [16], on autonomous system (AS) topology map information [5], or on third-party infrastructures like content delivery networks [6]. Motivated by the difficulty of inferring the ISPs' interests based on measurements, [3] and [4] investigated the use of ISP-provided information to influence peer selection. ISP-friendly peer selection needs careful design, as it may negatively affect the download performance and hence the user experience [7]. Furthermore, compared to P2P caching, which is controlled by the ISPs, ISP friendliness has to be implemented by the P2P protocol and is thus not directly controllable by the ISPs. Nonetheless, ISP-friendly peer selection is complementary to caching and can even benefit caching [14].

Caching of P2P contents has been the subject of several works. Most works focused on the achievable cache hit ratios [17], [18] and on the efficiency of various cache eviction policies [10], [11]. Our work is orthogonal to the works on cache eviction policies, as we assume the existence of a cache eviction policy, and we consider the impact of allocating the cache's upload bandwidth between competing overlays on the amount of inter-ISP traffic generated by the overlays.

Cache upload bandwidth management for P2P video streaming systems was considered in [19] and [20] in order to decrease the ISPs' incoming transit traffic. In the case of streaming, the download rate of peers is determined by the video rate, and the received rate does not influence the peers' behavior. This makes the problem of cache bandwidth allocation for streaming systems significantly different from the problem considered in this paper. We do not only consider the impact of the cache upload rate on the instantaneous inter-ISP traffic, but also its impact on the system dynamic.

Close to our work is [21], where the authors studied the impact of different bandwidth reservation schemes between two overlays via simulations. They concluded that the impact of cache bandwidth allocation was minor, which can be attributed to the inefficiency of the cache bandwidth utilization under the considered schemes. Compared to [21], in this paper we give a mathematical formulation of the problem of cache bandwidth allocation, use analytical models of the swarm dynamics and the inter-ISP traffic to give insight into the characteristics of an optimal allocation policy, and use simulations and experiments to demonstrate the inter-ISP traffic savings achievable through cache bandwidth allocation.

In [22]–[24], the authors proposed schemes for bandwidth allocation among multiple swarms in P2P systems. In these works, the initial seeder is the bandwidth allocator that attempts to maximize the total download rate of the system so as to minimize the download latencies of the peers. The most fundamental difference between our work and [22]–[24] is that we aim at minimizing the amount of inter-ISP traffic generated by the overlays. References [22] and [23] consider managed swarms, while in our work, caching is performed transparently to the peers. In [23], the authors assume peers belonging to multiple swarms. The bandwidth allocation among swarms is implemented by these peers as they follow the prioritization scheme suggested by the coordinator. More similar to our work is [24], where the authors implement a simple model-based controller for server bandwidth in BitTorrent systems. Due to the large amount of data needed to parametrize the model, the authors question the practicality of their approach [24].

Our work relies on the analytical models of the system dynamics of BitTorrent-like systems in [13] and [25]–[29]. These works used a Markovian model of the system dynamics of BitTorrent-like systems to model the service capacity and the scalability [25], [26], to evaluate the impact of peer upload rate allocation between two classes of peers [27], to assist the dimensioning of server assisted hybrid P2P content distribution [29], and to evaluate the impact of caches on the swarm dynamics and on the amount of inter-ISP traffic for a single overlay [13]. Our work differs significantly from these works, as we consider multiple overlays and use the fluid model of the system dynamics to get insight into the characteristics of an optimal P2P cache bandwidth allocation policy.

In our work, we model the cache bandwidth allocation problem in P2P systems as a Markov decision process (MDP). MDPs were used in [30]–[32] to analyze schemes for incentivizing fair resource reciprocation and for discouraging free-riding in P2P systems. Compared to [31]–[33], in our work we use an MDP to prove the existence of an optimal cache allocation policy.

III. SYSTEM MODEL

In the following, we describe our model of a multiswarm file-sharing system and our model of cache bandwidth allocation.

The model captures the effect of the cache bandwidth allocation on the evolution of the system.

We consider a set $\mathcal{I} = \{1, \ldots, I\}$ of ISPs, and a set of swarms $\mathcal{S} = \{1, \ldots, S\}$, whose peers are spread over the ISPs. Peers are either leechers, which download and upload simultaneously, or seeds, which upload only. Leechers arrive to swarm s according to a Poisson process with intensity λ_s; the arrival rate of leechers in ISP i is $\lambda_{i,s}$. The Poisson process can be a reasonable approximation of the arrival process over short periods of time [34], even if the arrival rate of peers varies over the lifetime of a swarm. We model the leechers' impatience by the abort rate θ. A leecher departs at this rate before downloading the entire content. Seeds depart from the swarm at rate γ, so that a seed stays on average $1/\gamma$ time in the swarm. The upload rate of peers is denoted by μ, and their download rate by c. We focus on the case when $\mu < c$. For simplicity, we consider that all files have the same size, and thus, μ and c can be normalized by the file size. Finally, we assume that leechers can use a share η of their upload rate due to partial content availability. This model of swarm dynamics was used in [13], [25], [26], [28], and [29].

We denote by $X_{i,s}(t)$ the number of leechers in ISP i in swarm s at time t, and by $Y_{i,s}(t)$ the number of seeds in ISP i in swarm s at time t. $X_{i,s}(t)$ and $Y_{i,s}(t)$ take values in the countably infinite state space $\mathbb{N}_0$. As a shorthand, we introduce $Z_{i,s}(t) = (X_{i,s}(t), Y_{i,s}(t))$ and $Z_s(t) = (Z_{i,s}(t))_{i \in \mathcal{I}}$. Finally, we denote the state of the swarms by $Z(t) = (Z_s(t))_{s \in \mathcal{S}}$.

Seeds and leechers in ISP i can upload and download data to and from peers in any ISP $j \in \mathcal{I}$. We define the publicly available upload rate $u_{i,s}^{\mathrm{P}}(t)$ as the available upload rate located in ISP i that can be used by leechers of swarm s in any ISP. This quantity tantamounts the upload rate of the leechers and the seeds $u_{i,s}^{\mathrm{P}}(t) = \mu(\eta X_{i,s}(t) + Y_{i,s}(t))$. A leecher cannot download from itself, therefore the publicly available upload rate in ISP i to a local leecher of swarm s is $u_{i,s}^{\mathrm{PL}}(t) = \max[0, \mu(\eta(X_{i,s}(t) - 1) + Y_{i,s}(t))]$.

A. P2P Cache Bandwidth Allocation Policies

The ISPs, as they are located in the lower layers of the ISP hierarchy, are interested in decreasing the inter-ISP traffic generated by the peers. In order to decrease its inter-ISP traffic, ISP $i \in \mathcal{I}$ maintains a cache with upload *bandwidth capacity* $K_i < \infty$, which acts as an ISP managed super peer [8]. The abstraction of a P2P cache as a source of upload bandwidth is motivated by the fact that P2P caches are often priced by their maximum upload rates. Since every ISP's goal is to decrease its own incoming inter-ISP traffic, it is reasonable to assume that the cache operated by ISP i only serves leechers in ISP i.

ISP i can implement an *active* cache bandwidth allocation policy to control the amount of cache bandwidth $\kappa_{i,s}(t)$ available to leechers in ISP i belonging to swarm s. We denote the cache bandwidth allocation of ISP i at time t by the vector $\kappa_i(t) = (\kappa_{i,1}(t), \ldots, \kappa_{i,S}(t))$, and the set of feasible cache bandwidth allocations of ISP i by $\mathcal{K}_i = \{\kappa_i \mid \sum_{s \in \mathcal{S}} \kappa_{i,s} \leq K_i\} \subseteq [0, K_i]^{|\mathcal{S}|}$. We also make the reasonable assumption that $\kappa_{i,s}(t) > 0$ for a swarm s only if the corresponding file is at least partially cached at ISP i at time t.

Parameter	Definition
$\mathcal{I}, \mathcal{S}$	Set of ISPs and set of swarms, respectively
$\kappa_{i,s}$	Cache bandwidth allocation of ISP i to swarm s
$\lambda_{i,s}$	Arrival rate of leechers to swarm s in ISP i
θ	Abort rate of leechers
γ	Departure rate of seeds
η	Effectiveness of file sharing
μ, c	Peer upload and download capacity, respectively
$X_{i,s}(t)$	Number of leechers in ISP i in swarm s at time t
$Y_{i,s}(t)$	Number of seeds in ISP i in swarm s at time t
$u_{i,s}^{PL}(t)$	Upload rate in ISP i available to all leechers in swarm s

Given the set $\mathcal{K}_i$ of feasible cache bandwidth allocations for ISP i, a cache bandwidth allocation *policy* π defines $\kappa_i(t)$ as a function of the system's history up to time t, i.e., $(Z(u))_{u<t}$, and past cache allocations $(\kappa_i(u))_{u<t}$. We denote the set of all cache bandwidth allocation policies by Π.

B. Caching and System Dynamics

Consider a policy π implemented by ISP i. We model the evolution of the swarms' state by an $I \times S \times 2$-dimensional continuous-time Markov jump process $\mathcal{Z}^\pi = \{Z(t), t \geq 0\}$, which is a collection of S coupled $I \times 2$-dimensional continuous-time Markov jump processes $\mathcal{Z}_s^\pi = \{Z_s(t), t \geq 0\}$.

Consider now a swarm $s \in \mathcal{S}$ under policy π, and denote the transition intensity from state z_s to state z_s' by $q_{z_s,z_s'}^\pi$. Denote by e_i the I-dimensional vector whose ith component is 1. The transition intensities from state $z_s = (x_s, y_s)$ are $q_{z_s,(x_s+e_i,y_s)}^\pi = \lambda_{i,s}$ (leecher arrival), $q_{z_s,(x_s-e_i,y_s)}^\pi = \theta x_{i,s}$ (leecher abort), and $q_{z_s,(x_s,y_s-e_i)}^\pi = \gamma y_{i,s}$ (seed departure). The transition intensity to state $(x_s - e_i, y_s + e_i)$, called the download completion rate, is a function of the maximum download rate of the leechers, and the available upload rate to leechers in ISP i.

1) Case of No Cache: Without a cache ($K_i = 0$), the leechers in ISP i would get a share $x_{i,s}/\sum_i x_{i,s}$ of the total upload rate $u_s^{\mathrm{P}} = \sum_i u_{i,s}^{\mathrm{P}}$ [25], [26], [28], [29]. The download completion rate in this case can be expressed as

$$q_{(x_s,y_s),(x_s-e_i,y_s+e_i)}^\pi = \min\left(cx_{i,s}, u_s^{\mathrm{P}} x_{i,s} / \sum_i x_{i,s}\right). \quad (1)$$

We refer to the process defined this way as the *uncontrolled* stochastic process, and we denote it by $\mathcal{Z}$.

2) Case of Cache: Consider that the instantaneous cache bandwidth allocated to swarm s is $\kappa_{i,s}$. The cache bandwidth increases the available upload rate, so that the download completion rate becomes

$$q_{(x_s,y_s),(x_s-e_i,y_s+e_i)}^\pi = \min\left(cx_{i,s}, u_s^{\mathrm{P}} x_{i,s} / \sum_i x_{i,s} + \kappa_{i,s}\right). \quad (2)$$

Since the cache bandwidth allocation can influence the transition intensities of the stochastic process, we refer to $\mathcal{Z}^\pi$ as the *controlled* stochastic process.

Table I summarizes the notation used in the paper.

IV. OPTIMAL CACHE BANDWIDTH ALLOCATION PROBLEM AND STATIONARY POLICY

In this section, we formulate the optimal cache bandwidth allocation problem and show the existence of an optimal stationary policy.

The primary goal of ISP i when allocating cache bandwidth to swarm s is to decrease the inter-ISP traffic. Cache bandwidth allocation inherently affects the upload rate available to the leechers, and hence it can affect the evolution of the process $\mathcal{Z}_s^\pi$.

Let us denote by $I_{i,s}(Z_s(t), \kappa_{i,s}(t))$ the rate of the incoming inter-ISP traffic in ISP i due to swarm s as a function of the cache bandwidth $\kappa_{i,s}(t)$ allocated to swarm s by ISP i and the swarm's state $Z_s(t)$. $I_{i,s}(Z_s(t), \kappa_{i,s}(t))$ also depends on $\kappa_{j,s}(t)$ of ISPs $j \neq i$, but as we focus on the bandwidth allocation problem of ISP i, for simplicity we assume that $\kappa_{j,s}(t) = \kappa_{j,s}$ is constant.

We can express the expected amount of incoming inter-ISP traffic under policy $\pi \in \Pi$ from time $t = 0$ until time T as

$$C_i^\pi(z, T) = E_z^\pi \left[\int_0^T \sum_{s \in \mathcal{S}} I_{i,s}(Z_s(t), \kappa_{i,s}(t)) dt \right]$$

where E_z^π denotes the expectation under policy π with initial state $Z(0) = z$.

Given the set Π of feasible cache bandwidth allocation policies, we define the cache bandwidth allocation problem for ISP i as finding the cache bandwidth allocation policy $\pi^* \in \Pi$ that minimizes the average incoming inter-ISP traffic rate $C_i^\pi(z)$ due to P2P content distribution, that is

$$\inf_{\pi \in \Pi} C_i^\pi(z) = \inf_\pi \limsup_{T \to \infty} \frac{1}{T} C_i^\pi(z, T). \quad (3)$$

Consequently, the optimal cache bandwidth allocation problem can be modeled as a continuous-time MDP with the optimality criterion defined in (3).

A. Optimal Cache Bandwidth Allocation

The first two fundamental questions that we are to answer are : 1) whether there is an optimal cache bandwidth allocation policy π^* that solves (3); and 2) whether there is an optimal policy whose choices only depend on the *current* system state $Z(t)$. Such a policy is called *stationary*. In general, an optimal stationary policy might not exist for an MDP when the action space or the state space is infinite. The following theorem shows that for the cache bandwidth allocation problem there exists an optimal stationary policy.

Theorem 1: There exists an optimal stationary policy π^* that minimizes the average traffic $C_i^\pi(z)$ of ISP i.

Proof of Theorem 1: Recall that the controlled processes $\mathcal{Z}_s^\pi$ are coupled through the bandwidth allocation policy π. In the following, we define four criteria C1–C4 for $\mathcal{Z}^\pi$ and use them to prove the theorem.

C1) The set $\mathcal{K}_i$ of cache bandwidth allocations is compact.
C2) For every state $z = (x, y)$, the incoming inter-ISP traffic rate $\sum_s I_{i,s}(z_s, \kappa_{i,s})$ and the transition intensities $(q_{(x_s, y_s),(x_s - e_i, y_s + e_i)}^\pi)_{s \in \mathcal{S}}$ are continuous functions of $\kappa_{i,s}$.

C3) Define $H(z) = C_i^\pi(z) - C_i^\pi(a)$, where a is an arbitrarily chosen state. Then, $\sum_{z'} H(z') q_{z,z'}^\pi$ is continuous in $\kappa_{i,s}$ for every state z.
C4) The average inter-ISP traffic $C_i^\pi(z)$ is finite for every policy π and initial state z.

We now formulate the following lemma based on [35, Theorem 5.9].

Lemma 1: For a continuous-time MDP with countably infinite state space and nonnegative cost, under C1–C4, there exists a stationary policy π^* that is average cost optimal.

Since the cost function $C_i^\pi(z)$ defined in (3) is the average cost, in order to prove the theorem, it is sufficient to show that $\mathcal{Z}^\pi$ fullfills the criteria C1–C4.

Proof of C1–C3: C1 follows from $0 \leq \kappa_{i,s}(t) \leq K_i < \infty$. $\sum_s I_{i,s}(z_s, \kappa_{i,s})$ is continuous by assumption, the continuity of the transition intensities $(q_{(x_s, y_s),(x_s - e_i, y_s + e_i)}^\pi)_{s \in \mathcal{S}}$ w.r.t. $\kappa_{i,s}$ follows from (2). C3 follows from the finiteness of $C_i^\pi(z)$ and from C2.

Proof of C4: In order to show the finiteness of the average inter-ISP traffic $C_i^\pi(z)$ for every policy π and initial state z, we show that $\mathcal{Z}_s^\pi$ satisfies the Foster–Lyapunov condition for every $s \in \mathcal{S}$, then we give a bound on the inter-ISP traffic rate in every state of the system. Let us define the Lyapunov function $w(z_s) = \sum_i (x_{i,s} + y_{i,s}) + 1$. Also, let us define the sequence $(t_n)_{n \geq 0}$ of time instants, which consists of the transition epochs of the process and of the instants when $\kappa_{i,s}(t)$ changes according to the policy π. Finally, we define the generalized average drift

$$AW(z_s) = E[w(Z_s(t_{n+1})) - w(Z_s(t_n)) | Z_s(t_n) = z_s]. \quad (4)$$

Consider now the Foster–Lyapunov average drift condition [36]

$$|AW(z_s)| < \infty \quad \forall z_s \quad \text{and} \quad AW(z_s) < -\epsilon \quad z_s \notin C \quad (5)$$

where $\epsilon > 0$ and $C \subset \mathbb{N}_0^{|\mathcal{I}| \times 2}$ is finite. For $\lambda_s < \infty$, the uncontrolled process $\mathcal{Z}_s$ satisfies (5): $|AW(z_s)| \leq 1$ due to the random-walk structure of the process, and $AW(z_s) = (\lambda_s - \theta x_s - \gamma y_s)/(-q_{z_s, z_s}) < -\epsilon$ for x_s or y_s sufficiently big. Consider now the mean drift $AW^\pi(z_s)$ of the controlled process. Again, $|AW^\pi(z_s)| \leq 1$. Furthermore, we have

$$AW^\pi(z_s) \leq AW(z_s) \frac{-q_{z_s, z_s}}{-q_{z_s, z_s} - K_i} < -\epsilon \frac{-q_{z_s, z_s}}{-q_{z_s, z_s} - K_i} < 0.$$

Consequently, the controlled process $\mathcal{Z}_s^\pi$ also satisfies the Foster–Lyapunov average drift condition. Since the process is aperiodic and irreducible, the drift condition guarantees ergodicity [36]. Furthermore, for $\tilde{M} = c > 0$, it holds that $I_{i,s}(z_s, \kappa_{i,s}) \leq \tilde{M} w(z_s)$. This together with the ergodicity of all $\mathcal{Z}_s^\pi$ implies that $C_i^\pi(z)$ is finite and concludes the proof. ∎

A consequence of Theorem 1 is that the optimal bandwidth allocation policy π^* is such that the allocation $\kappa_i(t)$ is only a function of the system state $Z(t)$, hence it is constant between the state transitions of $\mathcal{Z}^{\pi*}$.

The optimal policy π^* can be found using the policy iteration algorithm [35], but it requires the solution of the steady-state probabilities of the controlled Markov processes $\mathcal{Z}^\pi$. This can be prohibitive even for a moderate number of ISPs and swarms. In Section V, we propose and discuss different approximations that do not require finding the optimal stationary policy.

V. CACHE BANDWIDTH ALLOCATION POLICIES

In this section, we first discuss a baseline for bandwidth sharing. We then describe three approximations to the optimal cache bandwidth allocation policy.

Throughout the section, we assume that the inter-ISP traffic functions $I_{i,s}(z_s, \kappa_{i,s})$ are known and are continuous convex nonincreasing functions of $\kappa_{i,s}$. The assumptions of continuity, convexity, and nonincreasingness are rather natural.

A. Demand-Driven Bandwidth Sharing (DDS)

As a baseline for comparison, consider that ISP i does *not* actively allocate its cache bandwidth K_i, therefore leechers at different swarms compete with one another for cache bandwidth. The cache in ISP i maintains a drop-tail queue to store the requests received from the leechers in ISP i and serves the requests according to a first-in–first-out (FIFO) policy at the available upload bandwidth K_i. Let us denote by $\alpha_{i,s}$ the rate at which leechers of swarm s in ISP i request data from the cache in ISP i, and denote by $\sigma_{i,s}$ the mean service time of these requests. Then, the offered load of swarm s to the cache is $\rho_{i,s} = \alpha_{i,s}\sigma_{i,s}$. Clearly, if $\rho_{i,s} \geq 1$, then the FIFO queue is in a blocking state with probability $p_i^b > 0$.

If the requests from leechers in every swarm arrive according to a Poisson process, then the aggregate arrival process is Poisson. Since the arrival process is Poisson, an arbitrary request is blocked (i.e., dropped) with probability $p_{i,s}^b = p_i^b$ despite the possibly heterogeneous mean service times due to the PASTA property [37]. The effective (i.e., not blocked) load for swarm s can be expressed as $(1 - p_i^b)\rho_{i,s}$, and consequently the share of cache bandwidth used to serve requests for swarm s can be estimated as

$$\frac{\kappa_{i,s}}{\sum\limits_{s \in \mathcal{S}} \kappa_{i,s}} = \frac{(1 - p_i^b)\rho_s}{\sum\limits_{s \in \mathcal{S}} (1 - p_i^b)\rho_s} = \frac{\rho_s}{\sum\limits_{s \in \mathcal{S}} \rho_s}. \tag{6}$$

In general, if the arrival process of requests is not Poisson, then (6) does not hold. Nevertheless, as under the assumption of a Poisson request arrival process, the cache bandwidth is shared among the swarms proportional to the offered load (demand) of the swarms, we refer to this policy as the *demand-driven sharing* (DDS) policy.

B. One-Step Look Ahead (OLA) Allocation Policy

The one-step look ahead (OLA) policy π^{OLA} is a simple approximation of the optimal stationary cache bandwidth allocation policy π^*.

Consider the controlled Markov process $\mathcal{Z}^{\pi^{\mathrm{OLA}}}$, and let us denote the nth transition epoch of the process by t_n. Then, according to the OLA policy, the cache bandwidth allocation $\kappa_i(t)$ of ISP i for $t_n < t \leq t_{n+1}$ is such that it minimizes the incoming inter-ISP traffic rate given the state $Z(t_n) = z$ of the process $\mathcal{Z}^{\pi^{\mathrm{OLA}}}$

$$\kappa_i(t) = \arg\min_{\kappa_i \in \mathcal{K}_i} \sum_{s \in \mathcal{S}} I_{i,s}(z_s, \kappa_{i,s}). \tag{7}$$

By following the OLA policy, the ISP minimizes the incoming inter-ISP traffic in every state of the process $\mathcal{Z}^{\pi^{\mathrm{OLA}}}$. The OLA policy *adapts* to the system state, but unlike the optimal policy

π^*, it does not consider the impact of cache bandwidth allocation on the evolution of the number of peers.

Recall that, by assumption, $I_{i,s}(z_s, \kappa_{i,s})$ are continuous convex nonincreasing functions of $\kappa_{i,s}$ for every state z_s. In order to obtain the optimal solution to (7), consider the Lagrangian

$$L(z, \kappa_i, \zeta) = \sum_{s \in \mathcal{S}} I_{i,s}(z_s, \kappa_{i,s}) - \zeta \left(\sum_{s \in \mathcal{S}} \kappa_{i,s} - K_i \right) \tag{8}$$

where $\zeta \leq 0$ is the Lagrange multiplier. Then

$$\frac{\partial L(z, \kappa_i, \zeta)}{\partial \kappa_{i,s}} = \frac{\partial I_{i,s}(z_s, \kappa_{i,s})}{\partial \kappa_{i,s}} - \zeta \tag{9}$$

and

$$\frac{\partial L(z, \kappa_i, \zeta)}{\partial \zeta} = K_i - \sum_{s \in \mathcal{S}} \kappa_{i,s}. \tag{10}$$

Hence, a minimum of L over $\mathcal{K}_i$ is characterized by

$$\kappa_{i,s} > 0 \Rightarrow \frac{\partial_+ I_{i,s}(z_s, \kappa_{i,s})}{\partial \kappa_{i,s}} \geq \zeta \geq \frac{\partial_- I_{i,s}(z_s, \kappa_{i,s})}{\partial \kappa_{i,s}}$$

$$\kappa_{i,s} = 0 \Rightarrow \frac{\partial_- I_{i,s}(z_s, \kappa_{i,s})}{\partial \kappa_{i,s}} \geq \zeta$$

where ∂_+ and ∂_- denote the right and the left derivative of a semi-differentiable function. Since $\mathcal{K}_i$ is compact and convex, such a minimum exists and can be found using a projected subgradient method [38].

An important insight from the OLA policy is the following. If $I_{i,s}(z_s, \kappa_{i,s})$ are continuously differentiable, then at optimality, every swarm with nonzero cache bandwidth allocation provides equal marginal traffic saving. If $I_{i,s}(z_s, \kappa_{i,s})$ are not continuously differentiable, then for swarms with nonzero cache bandwidth allocation, the intersection of the subdifferentials is nonempty.

C. Steady-State Optimal (SSO) Allocation Policy

The opposite of the OLA policy is to focus on the long-term evolution of the controlled Markov process $\mathcal{Z}^\pi$, that is, on the incoming inter-ISP traffic in steady state, and to consider time-independent cache bandwidth allocation policies $\overline{\pi} = \kappa_i$.

Let us denote the expected number of leechers and seeds in steady state as a function of the cache bandwidth allocation policy $\overline{\pi}$ by $\overline{x}_{i,s}^{\overline{\pi}}$ and by $\overline{y}_{i,s}^{\overline{\pi}}$, respectively. They were shown to be a function of the cache upload rate $\kappa_{i,s}$ allocated to swarm s [13], [14]. As long as the total available upload rate is less than or equal to the total download rate of the leechers

$$\overline{x}_{i,s}^{\overline{\pi}} = \frac{\lambda_{i,s}}{\nu \left(1 + \frac{\theta}{\nu}\right)} - \frac{\kappa_{i,s}}{\mu\eta \left(1 + \frac{\theta}{\nu}\right)} - \Delta_i(\mathbf{x}, \mathbf{y}, \boldsymbol{\kappa}) \tag{11}$$

$$\overline{y}_{i,s}^{\overline{\pi}} = \frac{\lambda_{i,s}}{\gamma \left(1 + \frac{\theta}{\nu}\right)} + \frac{\kappa_{i,s}\theta}{\mu\eta\gamma \left(1 + \frac{\theta}{\nu}\right)} + \frac{\theta}{\gamma}\Delta_i(\mathbf{x}, \mathbf{y}, \boldsymbol{\kappa}) \tag{12}$$

where $\frac{1}{\nu} = \frac{1}{\eta}\left(\frac{1}{\mu} - \frac{1}{\gamma}\right) \geq 0$ [13], [26] and

$$\Delta_i(\mathbf{x}, \mathbf{y}, \boldsymbol{\kappa}) = \frac{\sum\limits_{j \in \mathcal{I}} (\lambda_{i,s}\kappa_{j,s} - \kappa_{i,s}\lambda_{j,s})}{\eta\gamma \left(1 + \frac{\theta}{\nu}\right) \left(\sum\limits_{j \in \mathcal{I}} (\lambda_{j,s} - \kappa_{j,s}) \right)}. \tag{13}$$

Otherwise, when the total upload rate exceeds the total download rate, increasing the cache bandwidth allocated to the swarm does not affect the number of leechers and seeds in steady state, which now depends on the peers' download capacity c [13], [26]

$$\overline{x}_{i,s}^{\overline{\pi}} = \frac{\lambda_{i,s}}{c\left(1 + \frac{\theta}{c}\right)} \quad \overline{y}_{i,s}^{\overline{\pi}} = \frac{\lambda_{i,s}}{\gamma\left(1 + \frac{\theta}{c}\right)}. \tag{14}$$

It is easy to verify that $\frac{\partial \overline{x}_{i,s}}{\partial \kappa_{i,s}} \leq 0$ and that $\frac{\partial^2 \overline{x}_{i,s}}{\partial \kappa_{i,s}^2} \geq 0$ for $\kappa_{i,s} \geq 0$, that is, the number of leechers in swarm s in ISP i in steady state is a convex nonincreasing function of the cache bandwidth allocated to swarm s in ISP i.

Given the functions $\overline{x}_{i,s}^{\pi}$ and $\overline{y}_{i,s}^{\pi}$ the steady-state optimal (SSO) bandwidth allocation policy can be formulated as

$$\overline{\pi}^* = \arg\min_{\kappa_i \in \mathcal{K}_i} \sum_{s \in \mathcal{S}} \overline{I}_{i,s}(\kappa_{i,s}) \tag{15}$$

where $\overline{I}_{i,s}(\kappa_{i,s})$ is the incoming inter-ISP traffic rate for the number of leechers and seeds in steady state.

Since by assumption $I_{i,s}(z_s, \kappa_{i,s})$ is convex nonincreasing in $\kappa_{i,s}$ for every state z_s, the steady-state optimal policy $\overline{\pi}^*$ can be found in a similar way as the *OLA* policy. The difference is that $\overline{I}_{i,s}(\kappa_{i,s})$ is a function of $\kappa_{i,s}$, $\overline{x}_{i,s}^{\pi}$ and $\overline{y}_{i,s}^{\pi}$, and the latter are themselves functions of $\kappa_{i,s}$. Note that the steady-state optimal policy $\overline{\pi}^*$ is not equivalent to the optimal policy π^* of the MDP, as the cache bandwidth allocated to a swarm s in ISP i would be nonzero even when $x_{i,s}(t) = 0$, which happens with nonzero probability.

D. Smallest-Ratio Priority Allocation

The *SSO* policy exclusively focuses on the long-term evolution of the process $\mathcal{Z}^\pi$ by minimizing the incoming inter-ISP traffic rate at steady state. It is time-independent, i.e., it does not adapt to the current state of the system. The *OLA* policy instead, adapts to the system state by minimizing the instantaneous incoming inter-ISP traffic rate. Nevertheless, the *OLA* policy disregards how cache bandwidth allocation affects the long-term evolution of the system.

In the following, we use the incoming inter-ISP traffic model in [14] to derive an adaptive cache bandwidth allocation policy that approximates the *SSO* policy.

1) Incoming Inter-ISP Traffic Model: We first reproduce the incoming inter-ISP traffic model for completeness. As the model is for a single swarm, we omit the subscript s for clarity. The model is based on two assumptions. First, leechers compete with each other for the available upload rate as long as they would be able to download at a higher rate. Second, given a single byte downloaded in ISP i, the distribution of its sources is proportional to the amount of upload rate exposed to the leechers that are located in ISP i.

The leechers in ISP i demand data at a total rate of cx_i. As the cache appears as an arbitrary peer to the leechers in ISP i, the demand is directed to the upload rate κ_i of ISP i's cache and to the publicly available upload rate $u_i^{\mathrm{PL}} + \sum_{j \neq i} u_j^{\mathrm{P}}$ of all ISPs. The leechers demand from the cache's upload rate with a probability proportional to its value, i.e., with probability $\kappa_i/(u_i^{\mathrm{PL}} + \sum_{j \neq i} u_j^{\mathrm{P}} + \kappa_i)$. The rest they demand from the publicly available upload rate, so the rate D_i^d that leechers in ISP

i demand from the publicly available upload rate can be expressed as

$$D_i^d = cx_i \left(1 - \frac{\kappa_i}{u_i^{\mathrm{PL}} + \sum_{j \neq i} u_j^{\mathrm{P}} + \kappa_i}\right). \tag{16}$$

If the system is limited by the download rate of the leechers, then the leechers receive the demanded rate. If the system is limited by the available upload rate, then the rate at which the leechers receive is proportional to the total publicly available upload rate divided by the total demanded rate

$$D_i^r = D_i^d \min\left(1, \frac{\sum_j u_j^{\mathrm{P}}}{\sum_j D_j^d}\right). \tag{17}$$

The rate that the leechers receive can originate from any ISP. Using the assumption that for a single byte downloaded in ISP i, the distribution of its sources is proportional to the amount of upload rate exposed to leechers in ISP i, we get the following estimate of the incoming inter-ISP traffic of ISP i:

$$I_i(z_s, \kappa_i) = D_i^r \left(\frac{\sum_{j \neq i} u_j^{\mathrm{P}}}{u_i^{\mathrm{PL}} + \sum_{j \neq i} u_j^{\mathrm{P}}}\right). \tag{18}$$

$I_i(z_s, \kappa_i)$ defined by (16)–(18) is a continuous convex nonincreasing function of the cache bandwidth κ_i allocated by ISP i.

2) Smallest-Ratio Priority Allocation: Our approximation of the *SSO* policy is based on the results in [13] and [14], which show that the dynamics of swarm s in ISP i only depend on the aggregate arrival intensity of leechers $\sum_{j \neq i} \lambda_{j,s}$ and on the aggregate cache capacity $\sum_{j \neq i} \kappa_{j,s}$ in the rest of the ISPs. This observation allows us to focus on a single swarm spread over two ISPs, $\mathcal{I} = \{1, 2\}$. ISP 1 is the tagged ISP, and ISP 2 is the aggregation of all other ISPs in the network. We denote the ratio of the arrival rates in the two ISPs by $r = \lambda_2/\lambda_1$.

Our focus will be on how the partial derivative $\frac{\partial \overline{I}_1(\kappa_1)}{\partial \kappa_1}$ of the steady-state inter-ISP traffic depends on r. For small κ_1, the incoming inter-ISP traffic $I_1(z, \kappa_1)$ of ISP 1 defined by (16) to (18) can be approximated by

$$I_1(z, \kappa_1) \approx \frac{x_1}{x_1 + x_2} u_2^{\mathrm{P}}. \tag{19}$$

We consider the case when the system is limited by the available upload rate, so we substitute (11) and (12) into (19) to obtain an approximation of the steady-state incoming inter-ISP traffic $\overline{I}_1(\kappa_1)$ of ISP 1 as a function of the cache bandwidth. Consider now the derivative at $\kappa_1 = 0$ and $\kappa_2 = 0$

$$\left.\frac{\partial \overline{I}_1(\kappa_1)}{\partial \kappa_1}\right|_{\substack{\kappa_1=0 \\ \kappa_2=0}} = -\frac{r^2(\gamma+\nu)(\gamma-\mu) - r\mu(\theta-\gamma)}{(1+\frac{\theta}{\nu})\mu\eta\gamma^2(1+r)^2}. \tag{20}$$

Recall that $\gamma - \mu > 0$ is a necessary condition for the upload rate to be the limit, and it implies $\nu > 0$ [13], [26]. Hence for $\theta - \gamma \leq 0$, (20) is negative and decreases monotonically in r.

For $\theta - \gamma > 0$, we have to consider the mixed second-order partial derivative at $\kappa_1 = 0$ and $\kappa_2 = 0$

$$\left.\frac{\partial^2 \overline{I}_1(\kappa_1)}{\partial \kappa_1 \partial r}\right|_{\substack{\kappa_1=0 \\ \kappa_2=0}} = -\frac{2r(\gamma+\nu)(\gamma-\mu) + (r-1)\mu(\theta-\gamma)}{(1+\frac{\theta}{\nu})\mu\eta\gamma^2(1+r)^3}. \tag{21}$$

Since $\theta - \gamma > 0$, (21) is negative for $r \geq 1$. Consequently, allocating cache bandwidth to swarms with a higher ratio r of arrival rates leads to a faster decrease of the steady-state inter-ISP traffic. At the same time, due to the term $(1+r)^3$ in the denominator $\lim_{r \to \infty} \left. \frac{\partial^2 \overline{I}_1(\kappa_1)}{\partial \kappa_1 \partial r} \right|_{\kappa_1 = 0} = 0$, i.e., swarms with a high arrival ratio r provide approximately the same gain.

This approximation suggests that a priority-based policy that assigns the highest priority to the swarms with highest ratio $r = \lambda_2/\lambda_1$ would resemble the *SSO* allocation policy for small cache bandwidths. We use this insight to define the *smallest-ratio priority* (*SRP*) cache bandwidth allocation policy. Under *SRP*, the priority of a swarm is calculated based on the instantaneous ratio of the local leechers to the number of peers in the overlay outside of ISP i, $\hat{r}_{i,s} = \frac{x_{i,s}(t)}{\sum_{j \neq i} z_{j,s}(t)}$. The priority of swarms with $\hat{r}_{i,s} = 0$ and $\hat{r}_{i,s} = \infty$ is lowest, and the priorities of the remaining swarms are assigned in decreasing order of the ratios $\hat{r}_{i,s}$. The priority is inverse proportional to $\hat{r}_{i,s}$ otherwise.

E. Practical Considerations

ISP i requires global information on the system state in order to compute any of the active cache bandwidth allocation policies presented above. The calculation of the *SSO* allocation relies on the incoming inter-ISP traffic rate for the number of peers in steady state. In order to compute (11) and (12), ISP i needs to estimate the arrival rates of leechers to the different swarms, i.e., $\lambda_{i,s} \, \forall s \in \mathcal{S}, \forall i \in \mathcal{I}$. The *OLA* policy assumes that the incoming inter-ISP traffic function $I_{i,s}(z_s, \kappa_{i,s})$ is known, furthermore it requires knowledge of the total number of peers in each swarm $z_s \, \forall s \in \mathcal{S}$. Similarly, under *SRP*, the priority of a swarm is calculated based on the number of local leechers $x_{i,s}$ and on the number of peers z_s in the overlay outside ISP i. In lack of interaction between ISPs and the P2P overlays, ISP i could potentially estimate the number of local peers in each swarm $z_{i,s}(t)$ by interrogating the local peers connected to the cache. Furthermore, ISP i can collect information about the aggregate number of peers outside its network $\sum_{j \neq i} z_{j,s}(t)$ by interrogating the tracker. If ISP-overlay interaction is possible [15], e.g., through an ALTO-like service [39], then ISP i could use the service to obtain information about the number of local and remote peers in each swarm. Cache bandwidth optimization could thus be a potential use-case for services like ALTO.

VI. PERFORMANCE EVALUATION AND INSIGHTS

In this section, we use simulations and experiments to compare the three approximate cache bandwidth allocation policies to *DDS*, and to provide insight into the characteristics of an optimal cache bandwidth allocation policy.

A. Performance Evaluation Setup

In the following, we describe the simulation and experimental settings and the implementation of the different cache bandwidth allocation policies. In both simulations and experiments, we consider $I = 2$ ISPs and use a BitTorrent seed with upload bandwidth K_1 as cache of ISP 1. The cache joins all swarms, but uploads only to leechers in ISP 1. The different cache bandwidth allocation policies are implemented in the peer.

1) Simulation Setup: We used the P2P simulation and prototyping toolkit ProtoPeer and the corresponding library for BitTorrent [40], [41] for the simulations. The simulations are flow-level: Data transmissions are flows, and the bandwidth for each flow is calculated according to the max-min-fair-share principle [42], an approximation of the bandwidth-sharing behavior of TCP.

We simulate 12–26 BitTorrent swarms, each sharing a file of 150 MB. The number of swarms is large enough to show the impact of the policies. At the same time, it keeps the run-time of the simulations at a reasonable level of a few hours per simulation run. The peers have an access bandwidth of 1 Mb/s upstream and 16 Mb/s downstream. The peers join swarm s in ISP i according to a Poisson process and, after completing the download, they remain in the swarm for an exponentially distributed seeding time with average $1/\gamma = 10$ min.

In order to implement the *OLA* and the *SSO* policies, we use the traffic model in Section V-D. We implement the *SRP* in the simulator by assigning a priority level to every data flow and by modifying the bandwidth-sharing algorithm. The bandwidth of flows with the same priority is calculated according to the original max-min-fair-share algorithm, while flows with a lower priority can only use the link bandwidth not used by flows with higher priority.

2) Experimental Setup: We perform experiments involving approximately 500 Planet-lab nodes using BitTorrent 4.4.0. We scale down the file size, the upload rates, and the download rates by a factor of 43 compared to the simulations in order to avoid interfering with other Planet-lab traffic: The file size is 3.5 MB, and the upload and download bandwidths of the peers are 23 and 373 kb/s, respectively.

For every swarm, we assign every Planet-lab node to one of the two ISPs and measure the traffic exchanged between peers belonging to different ISPs. We use one peer per swarm as the cache of ISP 1; these 12 peers run on a dedicated Linux computer. We implement the cache bandwidth allocation policies using hierarchical token bucket (HTB) queues in Linux traffic control. We use one filter per swarm to redirect the upload traffic of the 12 peers to an HTB class that enforces the total cache upload bandwidth limit K_1.

For the *SSO* and the *SRP* policies, we attach to this class one subclass per swarm. By default, each subclass has 500 B/s of guaranteed bandwidth in order to keep the TCP connections alive. The actual priority and guaranteed bit rate are then set according to the cache bandwidth allocation policy. The excess bandwidth is distributed among the swarms as defined by the HTB queue. For *SRP*, we update the priorities every 10 s based on the average number of leechers and seeds over the preceding 30 s.

B. Stationary Arrival Process

We start by considering the case when peers join swarm s in ISP i according to a stationary Poisson process at a rate of $\lambda_{i,s}$. This corresponds to a system in steady state.

Every simulation run corresponds to 6.5 h of simulated time, and we use the results following a warm-up period of 1.5 h. For every configuration, we show the average of five simulation runs together with the 95% confidence intervals. Every experiment

TABLE II
RELATIVE PEER ARRIVAL RATES IN THE SIMULATED SCENARIOS

Scenario	Number of swarms (S)	Identical swarms (s)	$\frac{\lambda_s}{\lambda}$	$\frac{\lambda_{2,s}}{\lambda_{1,s}}$
unif.,1:10	12	1,..,12	1/12	10
zipf,1:10	12		$\propto \frac{1}{s}$	10
unif.,1:1+1:10	12	1,..,10	1/12	10
		11,12	1/12	1
het.,2:2+1:10	15	1,..,4	1/8	10
		5,..,15	1/22	1

runs for 4 h, and we use the results after an initial warm-up period of 1 h.

1) Cache Bandwidth Allocation Matters: We simulate four scenarios to investigate under what conditions active cache bandwidth allocation can be beneficial. For simplicity, we denote the total arrival rate by $\lambda = \sum_i \sum_s \lambda_{i,s}$. We use the same total arrival rate $\lambda = 30$/min for all four scenarios, but the four scenarios differ in terms of the arrival rates $\lambda_{i,s}$ of the peers between swarms and between ISPs. Table II shows the relative arrival rates for the four scenarios. The ratio $\frac{\lambda_s}{\lambda}$ is related to the size of swarm s compared to all swarms, while $\frac{\lambda_{2,s}}{\lambda_{1,s}}$ is related to the share of local peers in swarm s. In [43], the authors measured the top-AS fraction of different swarms, defined as the maximum number of peers in one AS of the swarm normalized by the size of the swarm $\left(\max_i \frac{z_i}{\sum_j z_j} \right)$. The top-AS fraction was found to vary from a minimum of slightly less than 0.1, for swarms sharing international content, to a maximum of 0.5, for swarms sharing regional content. We can use these numbers to obtain estimates of the relative arrival intensities of leechers as follows. In absence of a cache ($\kappa_i = 0 \; \forall i \in \mathcal{I}$), the numbers of both seeders and leechers are proportional to the arrival rate of leechers to swarms, whether the system is upload-rate-limited (11)–(12) or download-rate-limited (14). Consequently, by substituting $\kappa_i = 0 \; \forall i \in \mathcal{I}$ in (11)–(14), we can use the top-AS fraction to calculate the ratio $\frac{\lambda_{2,s}}{\lambda_{1,s}}$. Using this approximation, a top-AS fraction of 0.1 corresponds to $\frac{\lambda_{2,s}}{\lambda_{1,s}}$ slightly greater than 9, and a top-AS fraction of 0.5 corresponds to $\frac{\lambda_{2,s}}{\lambda_{1,s}} = 1$. Given these approximate relative arrival intensities, our evaluation scenarios are constructed so that they allow us to isolate the factors that influence the efficiency of cache bandwidth allocation policies.

As an example, in scenario *unif.,1:1+1:10*, all $S = 12$ swarms have the same arrival rate $\lambda_s = \lambda/12$. The arrival rates for swarms 1–10 are asymmetric ($\lambda_{2,s} = 10\lambda_{1,s}$), while for swarms 11 and 12, they are symmetric ($\lambda_{2,s} = \lambda_{1,s}$). In scenario *het.,2:2+1:10*, the swarms have different arrival rates. Four out of 15 swarms have an arrival rate of $\lambda_s = \lambda/8$ and are asymmetric $\lambda_{2,s} = 10\lambda_{1,s}$. The remaining 11 swarms have an arrival rate of $\lambda_s = \lambda/22$ and are symmetric $\lambda_{2,s} = \lambda_{1,s}$. Compared to *unif.,1:1+1:10*, in this scenario the symmetric swarms, though more popular in ISP 1, are less popular in total than the asymmetric ones. The use of Zipf's law for the arrival intensities in scenario *zipf,1:10* is motivated by recent measurements that show that the distribution of the number of concurrent peers over swarms exhibits Zipf-like characteristics over a wide range of swarm sizes [44], [45]. Symmetric and asymmetric swarms are motivated by measurements that show
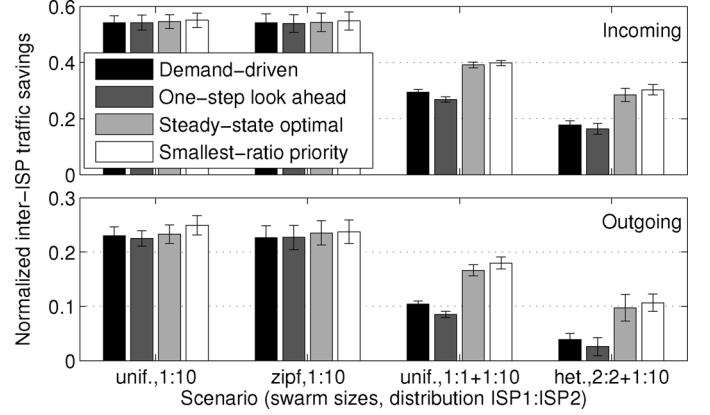

Fig. 1. Incoming and outgoing inter-ISP traffic savings for the four scenarios and four policies for $K_1 = 30$ Mb/s. Simulation results.

the difference in terms of the spatial distribution of peers between contents of regional and of global interest (e.g., the popularity of movies depending on the language [45]).

Fig. 1 shows the normalized incoming and outgoing inter-ISP traffic saving of ISP 1 for the four scenarios for the *DDS*, *OLA*, *SSO*, and *SRP* allocation policies. We calculate the normalized inter-ISP traffic saving as the decrease of the average inter-ISP traffic due to installing a cache divided by the average inter-ISP traffic without a cache ($K_1 = 0$), that is, $(C_i|_{K_1=0} - C_i^\pi)/C_i|_{K_1=0}$. The upload bandwidth of the cache in ISP 1 is $K_1 = 30$ Mb/s.

For the *unif.,1:10* and the *zipf,1:10* scenarios, in which the ratio $\lambda_{2,s}/\lambda_{1,s} = 10$ is the same for all swarms, the difference between the results for the different cache bandwidth allocation policies is within the confidence interval. However, for the scenarios *unif.,1:1+1:10* and *het.,2:2+1:10*, the bandwidth allocation policies make a significant difference in terms of traffic savings, both in terms of incoming and outgoing inter-ISP traffic. These results indicate that cache bandwidth allocation affects the inter-ISP traffic savings when the distribution of the peers over the ISPs is different among swarms, as for the *unif.,1:1+1:10* and the *het.,2:2+1:10* scenarios.

A comparison of the different policies in Fig. 1 for the *unif.,1:1+1:10* scenario reveals that the effect of the *OLA* policy on the inter-ISP traffic saving is opposite to the effect of the *SSO* and the *SRP* policies. The *OLA* policy performs worse than *DDS*, but the *SSO* and *SRP* policies compared to *DDS* drastically increase the incoming and outgoing inter-ISP traffic savings. For the *SSO* policy, the incoming inter-ISP traffic savings increase by about 33%, and the outgoing inter-ISP traffic savings by over 60%. For the *het.,2:2+1:10* scenario, the savings increase by 60% and 150%, respectively.

The *SRP* policy achieves even better gains. For the *het.,2:2+1:10* scenario for example, the savings increase by 71% for the incoming inter-ISP traffic and 172% for the outgoing traffic. Considering that P2P cache eviction policies were reported to achieve within 10%–20% of the hit rate of the optimal offline eviction policy [10], [11], i.e., very close to the cacheability of P2P content, the 30%–70% decrease of the incoming inter-ISP traffic achieved through cache bandwidth allocation is more than the potential decrease in cache miss
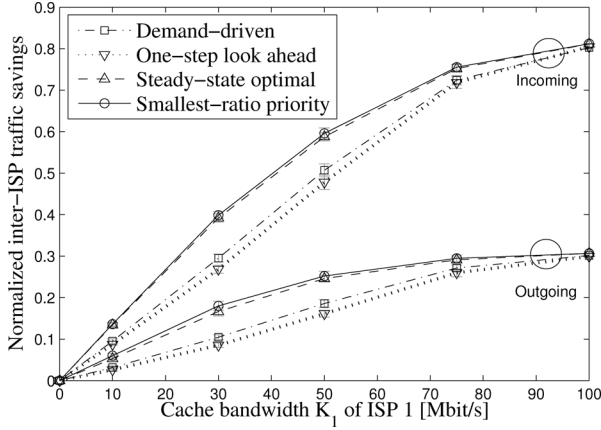
Fig. 2. Incoming and outgoing inter-ISP traffic saving for the *unif.,1:1+1:10* scenario versus cache bandwidth in ISP 1. Simulation results.

Fig. 3. Incoming and outgoing inter-ISP traffic saving for the *unif.,1:1+1:10* scenario versus cache bandwidth in ISP 1. Experiment results.

rate that could be achieved through improved cache eviction policies.

2) Inter-ISP Traffic Savings: Fig. 2 shows the incoming and outgoing inter-ISP traffic savings normalized by the inter-ISP traffic without cache ($K_1 = 0$) for the *unif.,1:1+1:10* scenario, as a function of the cache bandwidth K_1. The figure confirms that the observations made in Fig. 1 hold for a wide range of cache bandwidths K_1. Only above $K_1 \approx 75$ Mb/s, when the available upload bandwidth in ISP 1 exceeds the aggregate download bandwidth of the leechers within ISP 1, the marginal traffic saving diminishes and so does the difference between the policies. We note that the *SRP* policy performs slightly better than the *SSO* policy for all cache bandwidths. This is because the *SSO* allocation can be far from optimal when the instantaneous number of peers in the system is far away from the steady-state average number of peers. We show the corresponding experimental results in Fig. 3. We omit the results for the *OLA* policy since it performed poorly in all simulated scenarios. As shown in Fig. 3, the experimental results match the simulation results (cf. Fig. 2) and confirm the significant gain of cache bandwidth allocation observed in the simulations. The only difference is that the *SRP* policy performs slightly worse than in the simulations, which is due to the impact of the network-layer implementation of bandwidth allocation and priorities on TCP congestion control. An application-layer implementation of the policy could prevent this.

Fig. 4 shows the incoming and outgoing inter-ISP traffic saving normalized by the inter-ISP traffic without cache ($K_1 = 0$) for the *het.,2:2+1:10* scenario as a function of the cache bandwidth K_1. The figure allows us to draw similar conclusions as Fig. 2, except for the dip in the outgoing inter-ISP traffic saving for *DDS* at $K_1 = 10$ Mb/s. While surprising at first sight, the potential increase of the outgoing inter-ISP traffic due to caching for small, symmetric swarms (i.e., swarms 5–15) was pointed out in [13]. Since the *SRP* and the *SSO* policies allow little cache bandwidth to be used by the symmetric swarms for low K_1, they provide outgoing inter-ISP traffic savings even at $K_1 = 10$ Mb/s.

3) Cache Upload Rate to Swarms: In order to understand how the different policies allocate bandwidth to the different
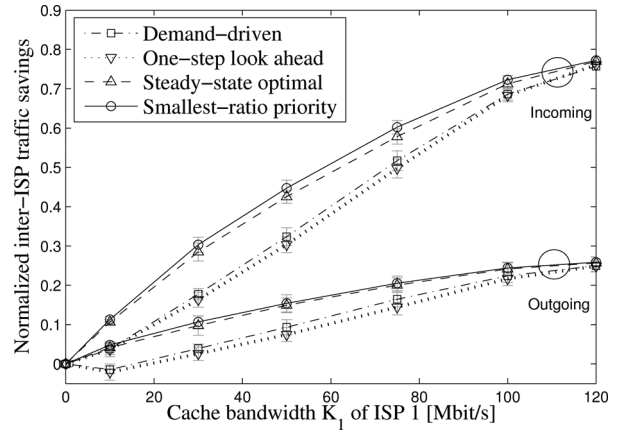

Fig. 4. Incoming and outgoing inter-ISP traffic savings for the *het.,2:2+1:10* scenario versus cache bandwidth in ISP 1. Simulation results.

swarms, we show two indifference maps of ISP 1 for the *unif.,1:1+1:10* scenario in Figs. 5 and 6. The horizontal and the vertical axes show the cache bandwidth allocated to each of the 10 asymmetric ($\lambda_{2,s} = 10\lambda_{1,s}$) and to each of the two symmetric ($\lambda_{2,s} = \lambda_{1,s}$) swarms, respectively. The curves show combinations of bandwidth allocations that lead to a particular inter-ISP traffic saving ΔI_1 (i.e., if there was no cache bandwidth constraint K_1, ISP 1 would be indifferent between allocations on the same indifference curve). The straight diagonal lines show different cache bandwidth constraints K_1. The *SSO* cache bandwidth allocation for K_1 is given by the coordinates of the point at which the cache bandwidth constraint line for K_1 is tangent to the indifference curve. The dotted line connects all such points: It shows the *SSO* cache bandwidth allocation for different K_1.

Fig. 5 shows the indifference map based on simulation results. We note that for $K_1 \leq 30$ Mb/s, all cache bandwidth should be allocated to the 10 asymmetric swarms, above that, as K_1 increases, so does the bandwidth that should be allocated to the two symmetric swarms. We also note that the shape of the indifference curves confirms that the inter-ISP traffic saving for a single swarm is a concave nondecreasing function of $\kappa_{i,s}$.

Fig. 6 shows the indifference map and the actual average cache upload rate received by the asymmetric (horizontal) and

Fig. 5.   Indifference map of ISP 1 for the *unif.,1:1+1:10* scenario based on simulation results. The dotted line shows the *SSO* cache bandwidth allocation for different values of cache bandwidth K_1.

Fig. 6.   Indifference map of ISP 1 for the *unif.,1:1+1:10* scenario based on the experiments, and the actual average cache upload rates for the *DDS*, *SSO*, and *SRP* policies. Experiment results.

the symmetric (vertical) swarms under the three allocation policies, based on experiment results. There is one marker per policy and total cache bandwidth K_1. The figure shows how the cache upload rate received by the swarms differs under the three policies depending on the cache bandwidth limit K_1. Under both *SRP* and *SSO*, the cache uploads to the symmetric swarms at a significantly lower rate than under *DDS* except for very high K_1, which is the key to the higher inter-ISP traffic savings of both policies.

C. Nonstationary Arrival Process

So far, we only looked at a system in steady state. In order to investigate the robustness of the bandwidth allocation policies to the system dynamics, we now turn to the case of a nonstationary arrival process. We consider that the leechers join swarm s in ISP i according to a nonstationary Poisson process with rate

$$\lambda_{i,s}(t) = \lambda_{i,s}^0 e^{-\frac{t}{\tau}} \quad (22)$$

where $\lambda_{i,s}^0$ is the initial arrival rate and τ is the attenuation parameter of peer arrival rate. Equation (22) has been shown to be a good model of the peer arrival rate during the swarm's lifespan in [46] and [47]. To derive the interarrival times of peers to swarm s, we simulated the nonstationary arrival process using the thinning method by Lewis and Schedler [48]. We considered $\sum_{i \in \mathcal{I}} \lambda_{i,s}^0 = \frac{1}{8}$ and $\tau = 4000$, which result in a swarm lifespan of about 5 h, and a total peer population during the lifespan of 500 peers. We performed experiments starting a new swarm every 15 min, for 6.5 h. The swarms starting at 0, 2, 4, and 6 h are symmetric ($\lambda_{2,s} = \lambda_{1,s}$), while the rest of the swarms are asymmetric ($\lambda_{2,s} = 10\lambda_{1,s}$).

Fig. 7 shows the incoming and outgoing inter-ISP traffic saving normalized by the inter-ISP traffic without cache ($K_1 = 0$) for the nonstationary scenario described above, as a function of the cache bandwidth K_1. The inter-ISP traffic savings show a similar trend as under the stationary arrival process (cf., Fig. 4). Comparing Figs. 7 and 4, we observe that the benefit of cache bandwidth allocation is slightly reduced, although still significant: The *SSO* and the *SRP* policies achieve savings in the order of 20%–30% compared to the *DD* allocation. It is
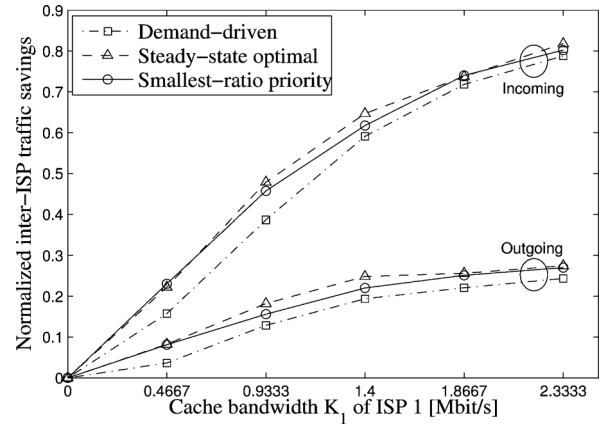

Fig. 7.   Incoming and outgoing inter-ISP traffic saving for the nonstationary scenario versus cache bandwidth in ISP 1. Experiment results.

important to note that the computation of the *SSO* policy and the derivation of the *SRP* policy in Section V-D assume that the system is in steady state, yet the policies provide significant savings in a nonstationary system.

VII. CONCLUSION

Motivated by the large amount of inter-ISP P2P traffic, we investigated a new dimension of P2P cache resource management, the allocation of cache upload bandwidth between overlays. We formulated the problem of cache bandwidth allocation as a Markov decision process and showed the existence of an optimal stationary allocation policy. Based on insights obtained from the model, we proposed three bandwidth allocation policies to approximate the optimal allocation policy. We performed simulations and experiments to evaluate the performance of the proposed policies. We demonstrated the importance of capturing the cache's impact on the swarm dynamics for cache bandwidth allocation. We identified the heterogeneity of the swarm's distribution between ISPs as the primary factor that influences the potential traffic savings through cache bandwidth allocation. Our results show that the proposed smallest ratio priority policy can decrease the amount of inter-ISP traffic between 30%–60%, which is higher than the potential decrease

in cache miss rate that could be achieved by improved P2P cache eviction policies given the cacheability of P2P traffic reported in the literature.

An interesting continuation of the work presented in the paper would be to evaluate the proposed cache bandwidth allocation policy together with other solutions. One could evaluate the savings that could be achieved by jointly allocating the cache bandwidth and the cache storage (through insertion and eviction policies). Similarly, it would be interesting to compare the proposed adaptive cache bandwidth allocation scheme with locality-aware neighbor-selection schemes, and to investigate the potential interaction between these solutions. We leave these interesting problems to be subject of future work.

REFERENCES

[1] H. Schulze and K. Mochalski, "Internet study 2008/2009," 2009 [Online]. Available: http://www.ipoque.com/resources/internet-studies
[2] P. Eckersley, F. von Lohmann, and S. Schoen, "Packet forgery by ISPs: A report on the Comcast affair," White Paper, Nov. 2007.
[3] V. Aggarwal, A. Feldmann, and C. Scheideler, "Can ISPs and P2P systems co-operate for improved performance?," *Comput. Commun. Rev.*, vol. 37, no. 3, pp. 29–40, Jul. 2007.
[4] H. Xie, Y. R. Yang, A. Krishnamurthy, Y. G. Liu, and A. Silberschatz, "P4P: Provider portal for applications," *Comput. Commun. Rev.*, vol. 38, no. 4, pp. 351–362, Oct. 2008.
[5] R. Bindal *et al.*, "Improving traffic locality in BitTorrent via biased neighbor selection," in *Proc. IEEE ICDCS*, Jul. 2006, pp. 66–75.
[6] D. R. Choffnes and F. E. Bustamante, "Taming the torrent: A practical approach to reducing cross-ISP traffic in peer-to-peer systems," *Comput. Commun. Rev.*, vol. 38, no. 4, pp. 363–374, Oct. 2008.
[7] S. L. Blond, A. Legout, and W. Dabbous, "Pushing BitTorrent locality to the limit," *Comput. Netw.*, vol. 55, no. 3, pp. 541–557, 2011.
[8] Allot Communications, Woburn, MA, USA, "Allot MediaSwift E," [Online]. Available: http://www.allot.com/MediaSwift_Video_Acceleration.html
[9] PeerApp, Ltd., Newton, MA, USA, "PeerApp UltraBand," [Online]. Available: http://www.peerapp.com
[10] A. Wierzbicki, N. Leibowitz, M. Ripeanu, and R. Woźniak, "Cache replacement policies for P2P file sharing protocols," *Euro. Trans. Telecommun.*, vol. 15, pp. 559–569, Nov. 2004.
[11] M. Hefeeda and O. Saleh, "Traffic modeling and proportional partial caching for peer-to-peer systems," *IEEE/ACM Trans. Netw.*, vol. 16, no. 6, pp. 1447–1460, Dec. 2008.
[12] N. Megiddo and D. Modha, "ARC: A self-tuning, low overhead replacement cache," in *Proc. USENIX FAST*, 2003, pp. 115–130.
[13] F. Lehrieder, G. Dán, T. Hoßfeld, S. Oechsner, and V. Singeorzan, "The impact of caching on BitTorrent-like peer-to-peer systems," in *Proc. IEEE P2P*, Aug. 2010, pp. 1–10.
[14] F. Lehrieder, G. Dan, T. Hossfeld, S. Oechsner, and V. Singeorzan, "Caching for bittorrent-like P2P systems: A simple fluid model and its implications," *IEEE/ACM Trans. Netw.*, vol. 20, no. 4, pp. 1176–1189, Aug. 2012.
[15] G. Dán *et al.*, "Interaction patterns between P2P content distribution systems and ISPs," *IEEE Commun. Mag.*, vol. 49, no. 5, pp. 222–230, May 2011.
[16] S. Ren *et al.*, "TopBT: A topology-aware and infrastructure-independent BitTorrent," in *Proc. IEEE INFOCOM*, Apr. 2011, pp. 1–9.
[17] N. Leibowitz, A. Bergman, R. Ben-Shaul, and A. Shavit, "Are file swapping networks cacheable? Characterizing P2P traffic," in *Proc. WCW*, Aug. 2002, pp. 1–6.
[18] T. Karagiannis, P. Rodriguez, and K. Papagiannaki, "Should Internet service providers fear peer-assisted content distribution?," in *Proc. ACM SIGCOMM IMC*, 2005, pp. 63–76.

[19] G. Dan, "Cooperative caching and relaying strategies for peer-to-peer content delivery," in *Proc. IPTPS*, 2008, pp. 1–16.
[20] J. Dai, B. Li, F. Liu, B. Li, and H. Jin, "On the efficiency of collaborative caching in ISP-aware P2P networks," in *Proc. IEEE INFOCOM*, 2011, pp. 1224–1232.
[21] I. Papafili, G. D. Stamoulis, F. Lehrieder, B. Kleine, and S. Oechsner, "Cache capacity allocation to overlay swarms," in *Proc. Int. Workshop Self-Org. Syst.*, Feb. 2011, pp. 68–80.
[22] R. S. Peterson and E. G. Sirer, "Antfarm: Efficient content distribution with managed swarms," in *Proc. NSDI*, 2009, pp. 107–122.
[23] R. S. Peterson, B. Wong, and E. G. Sirer, "A content propagation metric for efficient content distribution," in *Proc. ACM SIGCOMM*, 2011, pp. 326–337.
[24] A. Sharma, A. Venkataramani, and A. A. Rocha, "Pros & cons of model-based bandwidth control for client-assisted content delivery," in *Proc. COMSNETS*, 2014, pp. 1–8.
[25] X. Yang and G. de Veciana, "Service capacity of peer to peer networks," in *Proc. IEEE INFOCOM*, Mar. 2004, pp. 2242–2252.
[26] D. Qiu and R. Srikant, "Modeling and performance analysis of BitTorrent-like peer-to-peer networks," in *Proc. ACM SIGCOMM*, Aug. 2004, pp. 367–378.
[27] F. Clévenot-Perronnin, P. Nain, and K. W. Ross, "Multiclass P2P networks: Static resource allocation for service differentiation and bandwidth diversity," *Perform. Eval.*, vol. 62, pp. 32–49, Oct. 2005.
[28] Y. Tian, D. Wu, and K. W. Ng, "Modeling, analysis and improvement for BitTorrent-like file sharing networks," in *Proc. IEEE INFOCOM*, Apr. 2006, pp. 1–11.
[29] I. Rimac, A. Elwalid, and S. Borst, "On server dimensioning for hybrid P2P content distribution networks," in *Proc. IEEE P2P*, Sep. 2008, pp. 321–330.
[30] R. Izhak-Ratzin, H. Park, and M. van der Schaar, "Online learning in BitTorrent systems," *IEEE Trans. Parallel Distrib. Syst.*, vol. 23, no. 12, pp. 2280–2288, Dec. 2012.
[31] E. J. Friedman, J. Y. Halpern, and I. Kash, "Efficiency and Nash equilibria in a scrip system for P2P networks," in *Proc. ACM EC*, 2006, pp. 140–149.
[32] H. Park and M. V. D. Schaar, "A framework for foresighted resource reciprocation in P2P networks," *IEEE Trans. Multimedia*, vol. 11, no. 1, pp. 101–116, Jan. 2009.
[33] R. Izhak-Ratzin, H. Park, and M. van der Schaar, "Reinforcement learning in BitTorrent systems," in *Proc. IEEE INFOCOM*, 2011, pp. 406–410.
[34] L. Guo *et al.*, "Measurement, analysis, and modeling of BitTorrent-like systems," in *Proc. ACM IMC*, Oct. 2005, pp. 35–48.
[35] X. Guo and O. Hernandez-Lerma, *Continuous-time Markov Decision Processes*, ser. Stochastic Modelling and Applied Probability. New York, NY, USA: Springer, 2009.
[36] R. L. Tweedie, "Criteria for ergodicity, exponential ergodicity and strong ergodicity of Markov processes," *J. Appl. Probab.*, vol. 18, pp. 122–130, 1981.
[37] R. W. Wolff, "Poisson arrivals see time averages," *Oper. Res.*, vol. 30, no. 2, pp. 223–231, 1982.
[38] N. Z. Shor, *Minimization Methods for Non-Differentiable Functions*, ser. Computational Mathematics. New York, NY, USA: Springer, 1985.
[39] M. Stiemerling, S. Kiesel, S. Previdi, and M. Scharf, "ALTO deployment considerations," Internet Engineering Task Force (IETF), Internet-Draft, 2013.
[40] "ProtoPeer," [Online]. Available: http://sourceforge.net/projects/protopeer/
[41] W. Galuba, K. Aberer, Z. Despotovic, and W. Kellerer, "ProtoPeer: A P2P toolkit bridging the gap between simulation and live deployment," in *Proc. Int. Conf. Simulat. Tools Tech.*, 2009, pp. 1–9.
[42] D. Bertsekas and R. Gallagher, *Data Networks*. Upper Saddle River, NJ, USA: Prentice-Hall, 1987.
[43] T. H. feld *et al.*, "Characterization of BitTorrent swarms and their distribution in the Internet," *Comput. Netw.*, vol. 55, no. 5, pp. 1197–1215, 2011.
[44] G. Dán and N. Carlsson, "Power-law revisited: A large scale measurement study of P2P content popularity," in *Proc. IPTPS*, Apr. 2010, pp. 1–5.
[45] T. Hoßfeld *et al.*, "Characterization of BitTorrent swarms and their distribution in the Internet," *Comput. Netw.*, vol. 55, no. 5, pp. 1197–1215, Apr. 2011.
[46] L. Guo *et al.*, "A performance study of BitTorrent-like peer-to-peer systems," *IEEE J. Sel. Areas Commun.*, vol. 25, no. 1, pp. 155–169, Jan. 2007.

[47] N. Carlsson, G. Dán, A. Mahanti, and M. Arlitt, "A longitudinal characterization of local and global BitTorrent workload dynamics," in *Proc. PAM*, 2012, pp. 252–262.

[48] P. Lewis and G. Shedler, "Simulation of nonhomogeneous Poisson processes by thinning," *Nav. Res. Logistics Quart.*, vol. 26, pp. 403–413, 1979.

Frank Lehrieder studied computer science and business administration at the University of Würzburg, Würzburg, Germany, and the University Antonio de Nebrija, Madrid, Spain. He received the Master's and Ph.D. degrees in computer science from the University of Würzburg in 2008 and 2013, respectively.

His special interests are performance evaluation and resource management for peer-to-peer networks and Future Internet infrastructures.

Valentino Pacifici studied computer engineering at the Politecnico di Milano, Milan, Italy. He received the joint M.Sc. degree in computer engineering from the Politecnico di Milano and KTH, Royal Institute of Technology, Stockholm, Sweden, in 2010, and is currently pursuing the Ph.D. degree in electrical engineering at KTH, Royal Institute of Technology.

He is a now a Researcher with the Laboratory for Communication Networks, KTH, Royal Institute of Technology. His research interests include the modeling, design, and analysis of content management systems using game-theoretical and stochastic tools.

György Dán received the M.Sc. degree in computer engineering from the Budapest University of Technology and Economics, Budapest, Hungary, in 1999, the M.Sc. degree in business administration from the Corvinus University of Budapest, Budapest, Hungary, in 2003, and the Ph.D. degree in telecommunications from KTH, Royal Institute of Technology, Stockholm, Sweden, in 2006.

He is an Associate Professor with KTH, Royal Institute of Technology. He worked as a consultant in the field of access networks, streaming media, and video conferencing from 1999 to 2001. He was a Visiting Researcher with the Swedish Institute of Computer Science, Kista, Sweden, in 2008, and a Fulbright Research Scholar with the University of Illinois at Urbana-Champaign, Urbana, IL, USA, from 2012 to 2013. His research interests include the design and analysis of content management and computing systems, game-theoretical models of networked systems, and cyber-physical system security in power systems.